

Cloud Computing Applications and Services

(Aplicações e Serviços de Computação em Nuvem)

Virtualization (Part II)

University of Minho

2024-2025



Containers

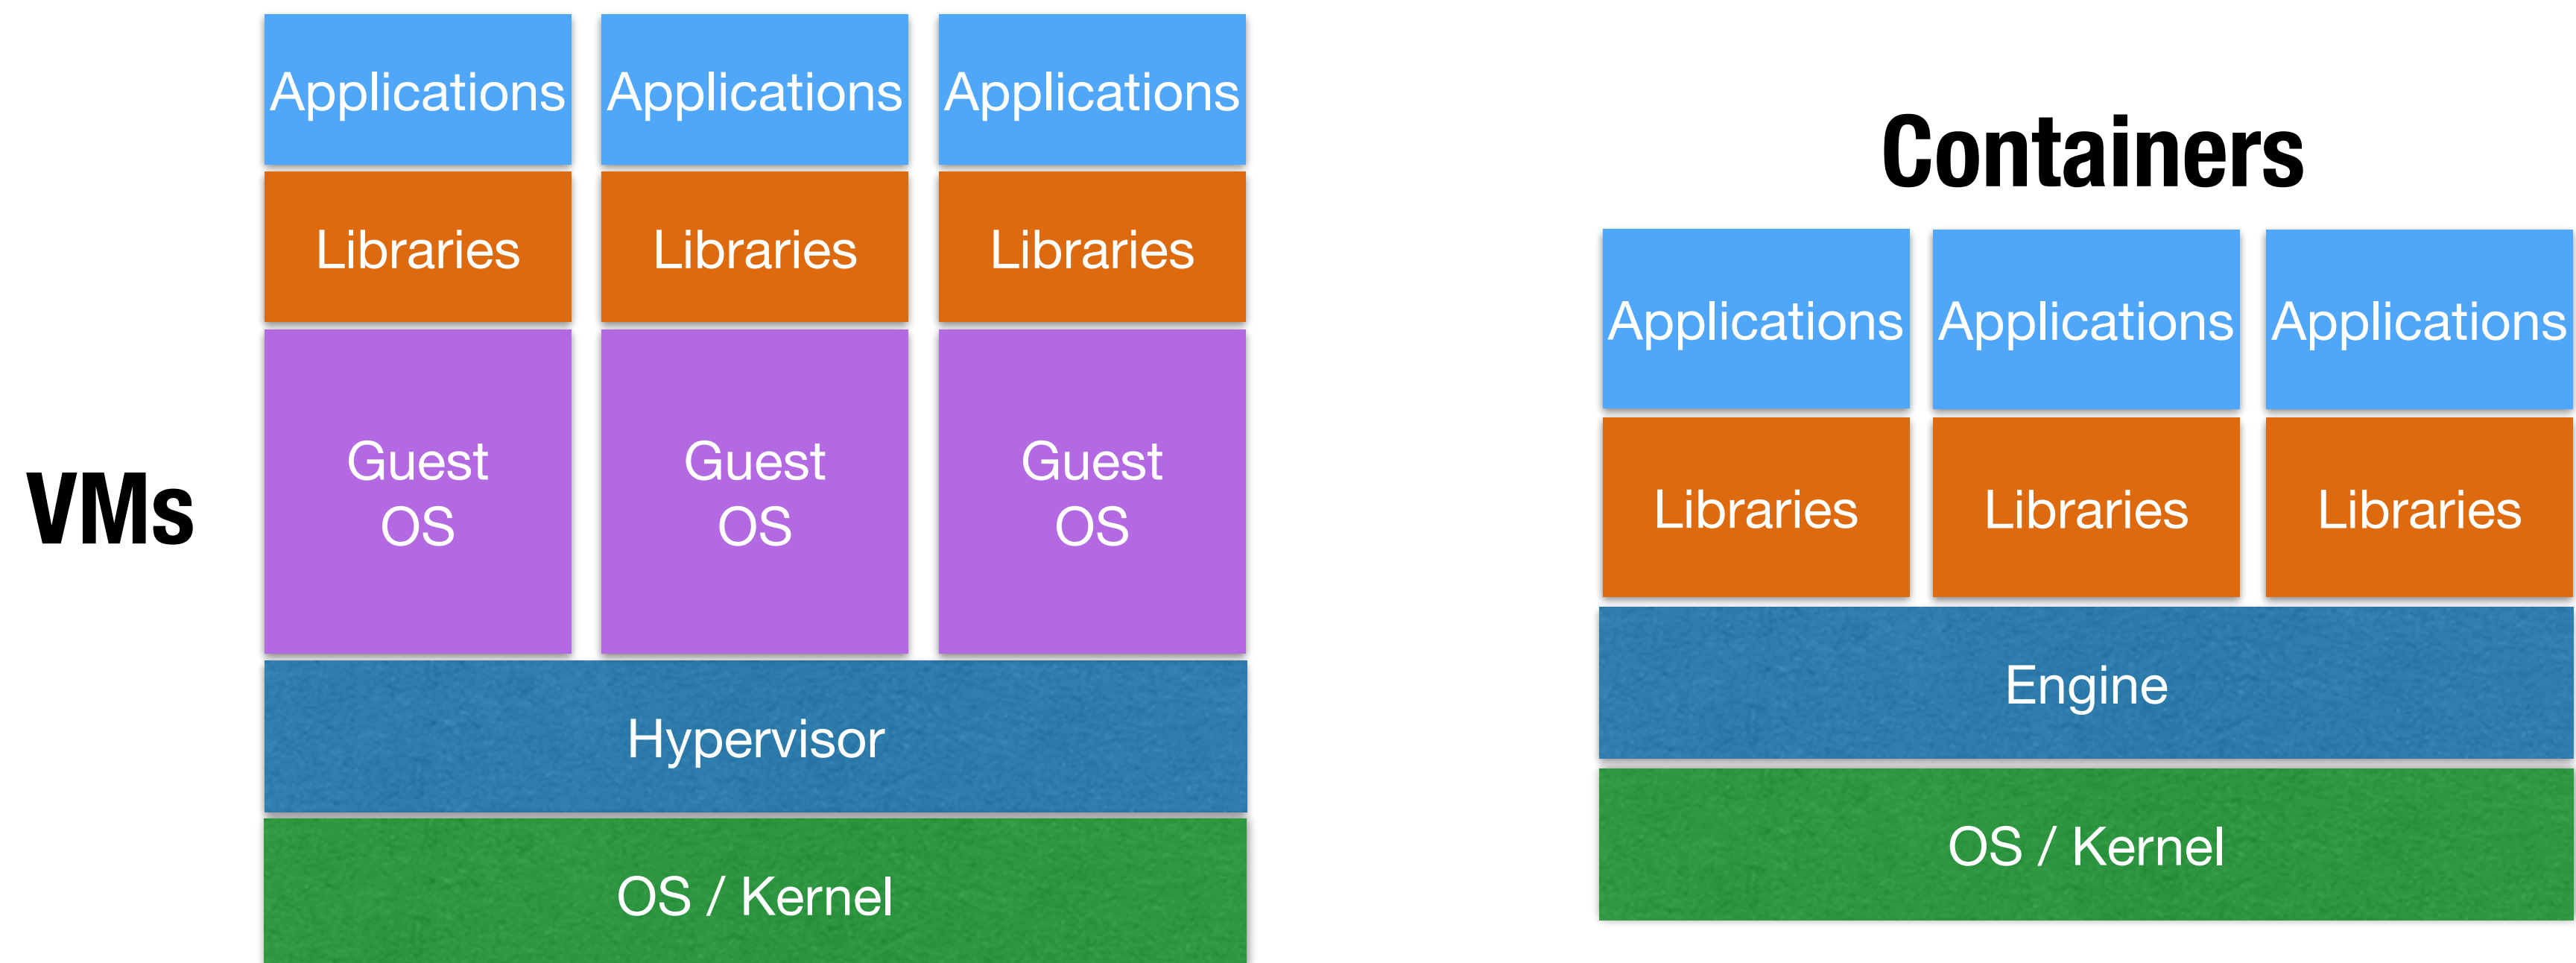


- **Lightweight virtual environment** that **groups and isolates a set of processes and resources** (memory, CPU, disk, ...), from the host and any other containers
- Why are containers useful?
 - Running different isolated versions of the same software/application (e.g., database) in a shared OS/Kernel environment
 - Portability/migration across servers
 - Easy packaging of software, applications and their dependencies

Linux Containers

Architecture

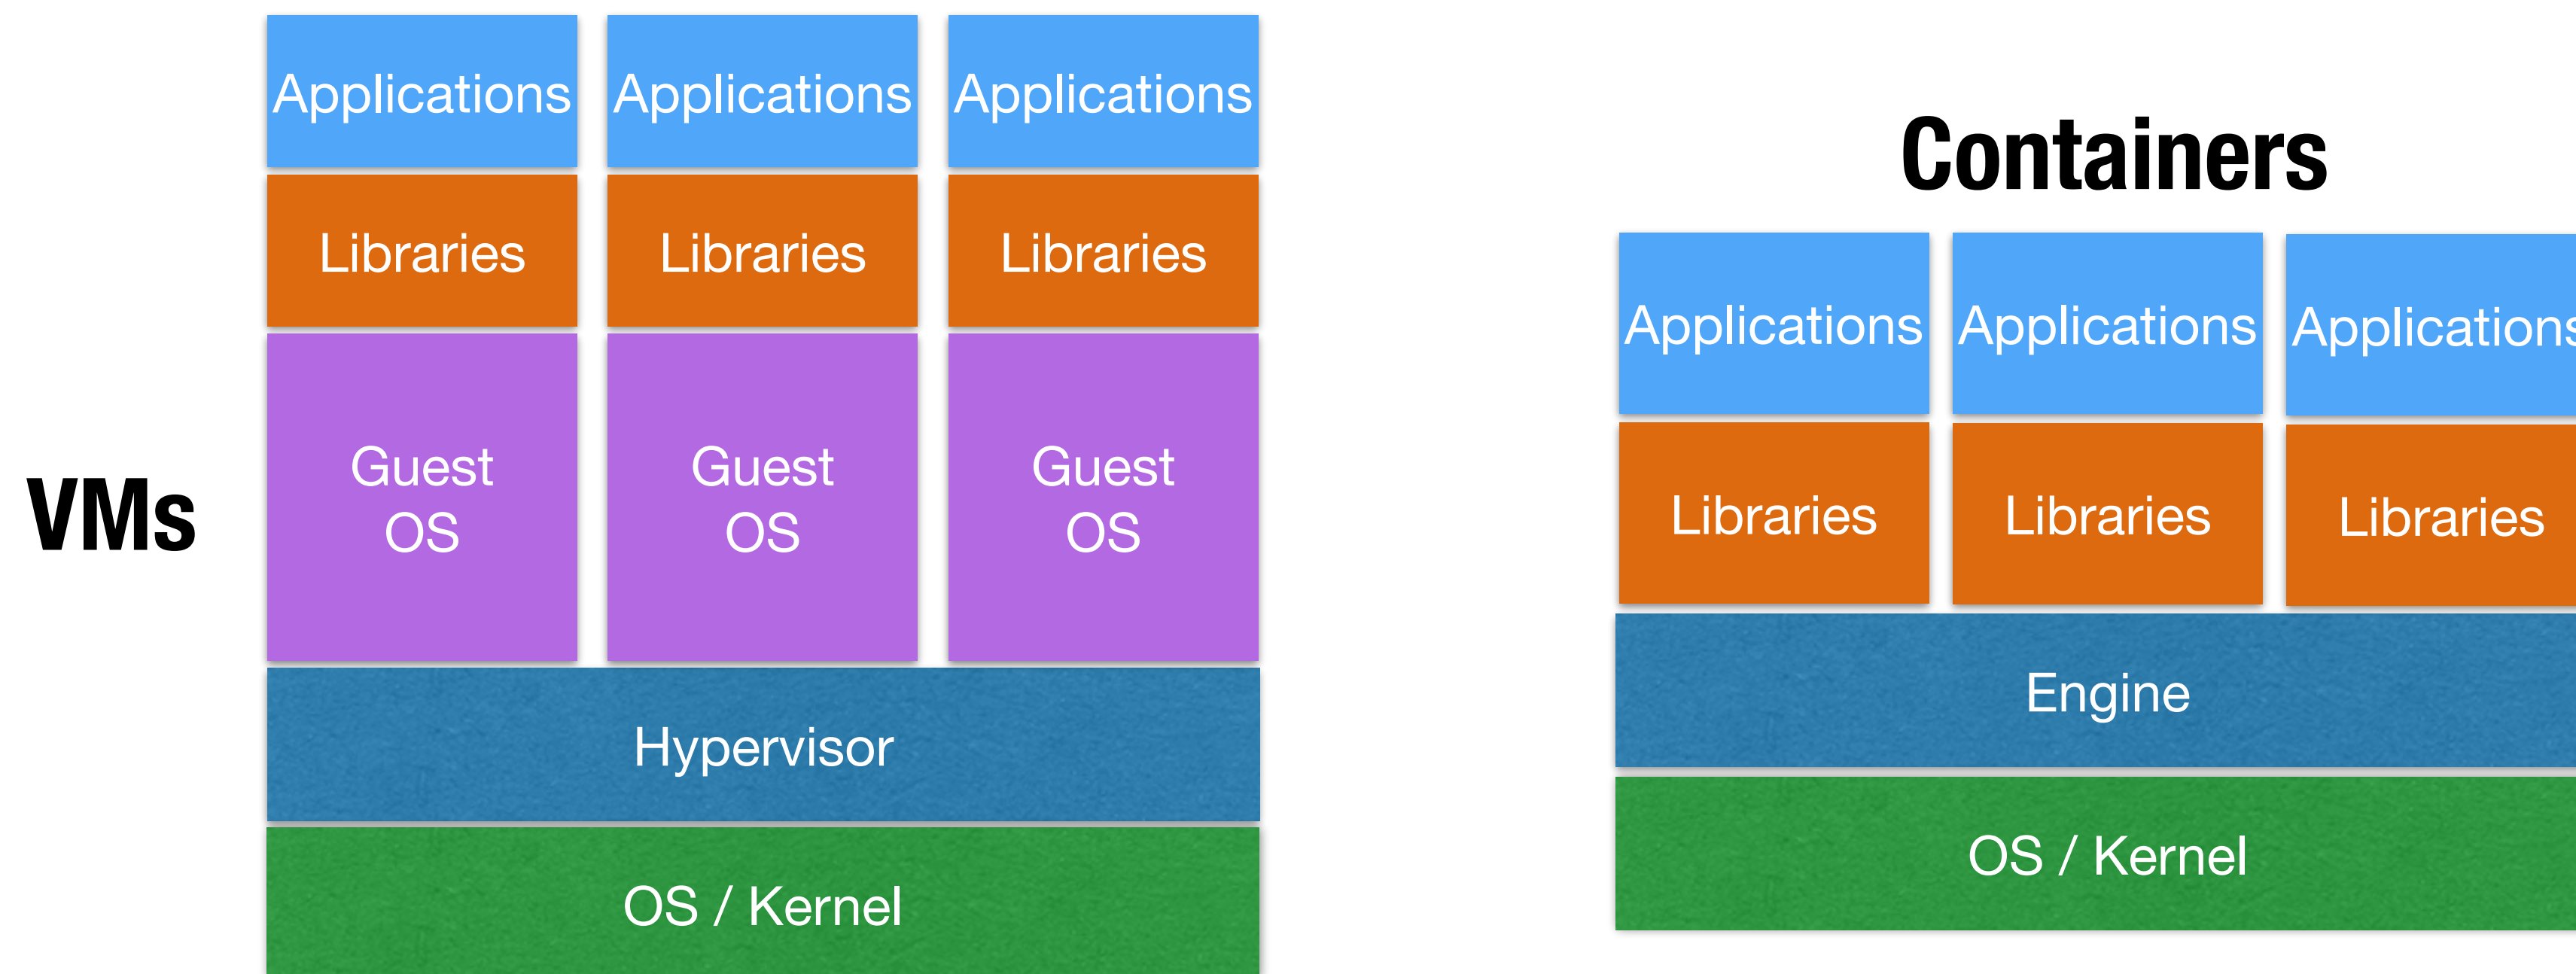
- Containers are lightweight, while **not requiring full virtualization** of CPU, RAM, network, and storage requests (as in VMs)



Linux Containers

Engine

- The **engine** isolates and configures resources
 - Containers are isolated from each other, i.e., the host is compartmentalized in terms CPU, RAM, memory, disk
 - Each container shares the hardware and kernel/OS with the host system



Linux Containers

Building Blocks (components of Linux kernel)

● Namespaces (Isolation)

- Host resources (e.g., network, filesystem) are **partitioned** into dedicated resources that are only accessible by a certain group of processes (under the same namespace)
- Allow **isolating** host resources across different containers

● Control Groups (Resource Management)

- Allow **allocating resources** (CPU, RAM, Disk, network) among groups of processes
- **Limit the amount of resources** per container (e.g., CPU cores, RAM/storage)

● SELinux (Security)

- Provides **additional security** over namespaces so that a container is not able to compromise the host system and other containers
- Enforces **access control and security policies**

Linux Containers

Types

- **OS Containers** (*e.g.*, LXC)

- Containers run multiple processes and simulate a “lightweight” operating system

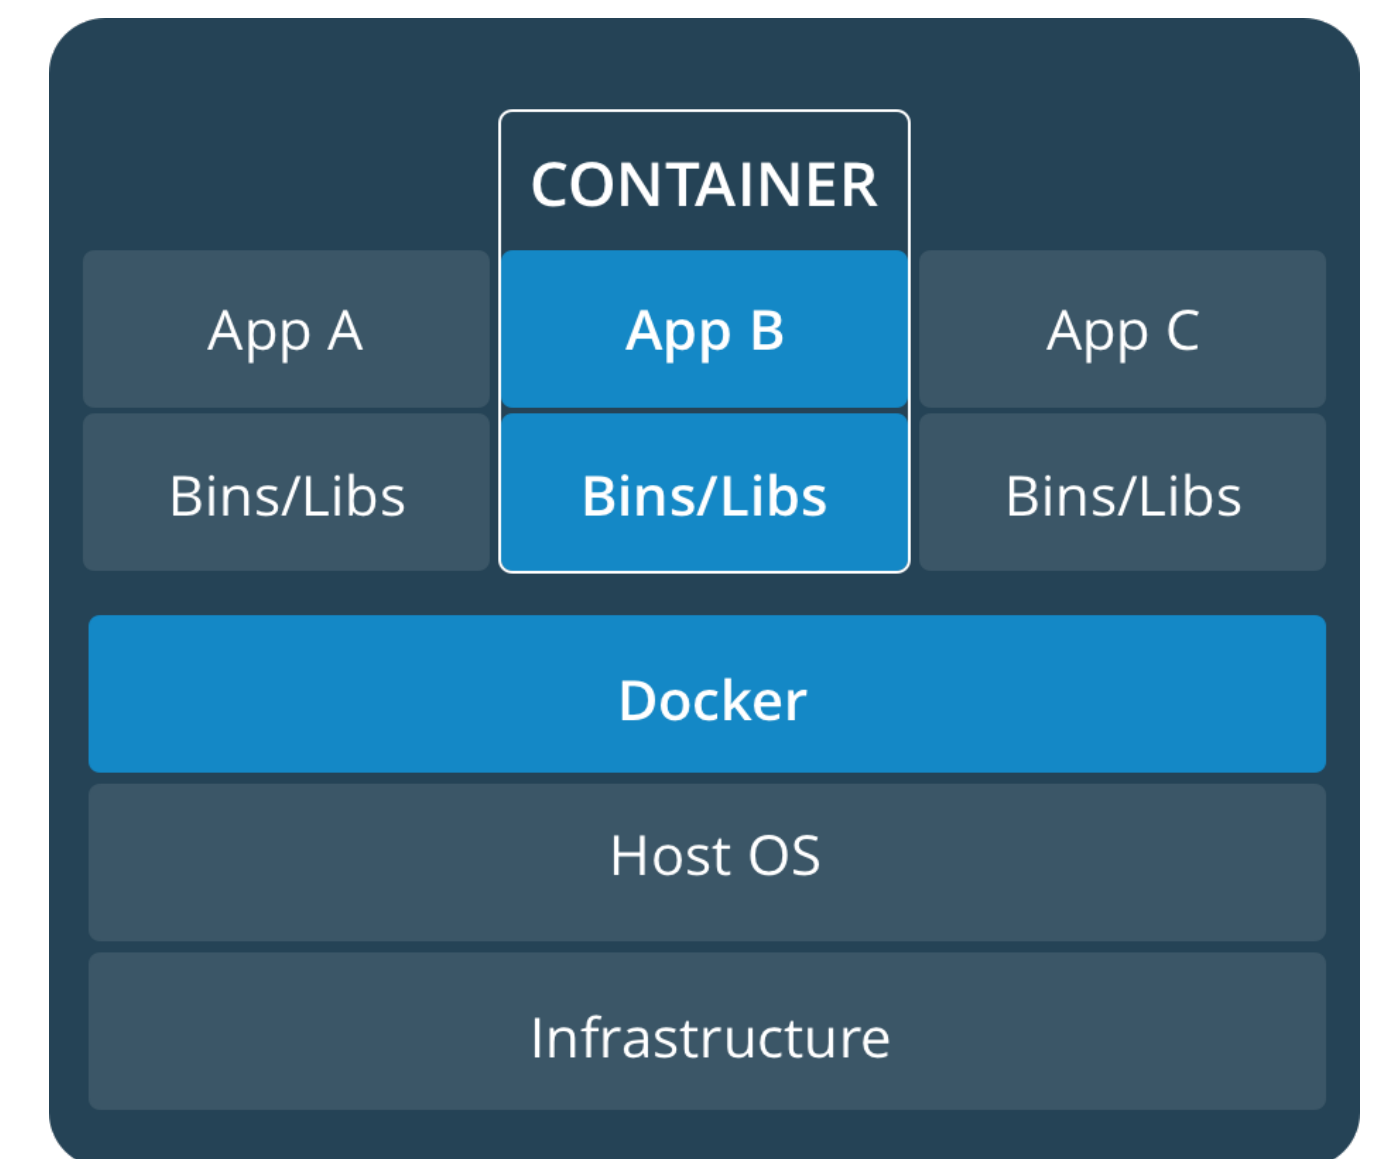
- **Application Containers** (*e.g.*, Docker)

- Focused on deploying applications
- Each application is seen as an independent process

Docker



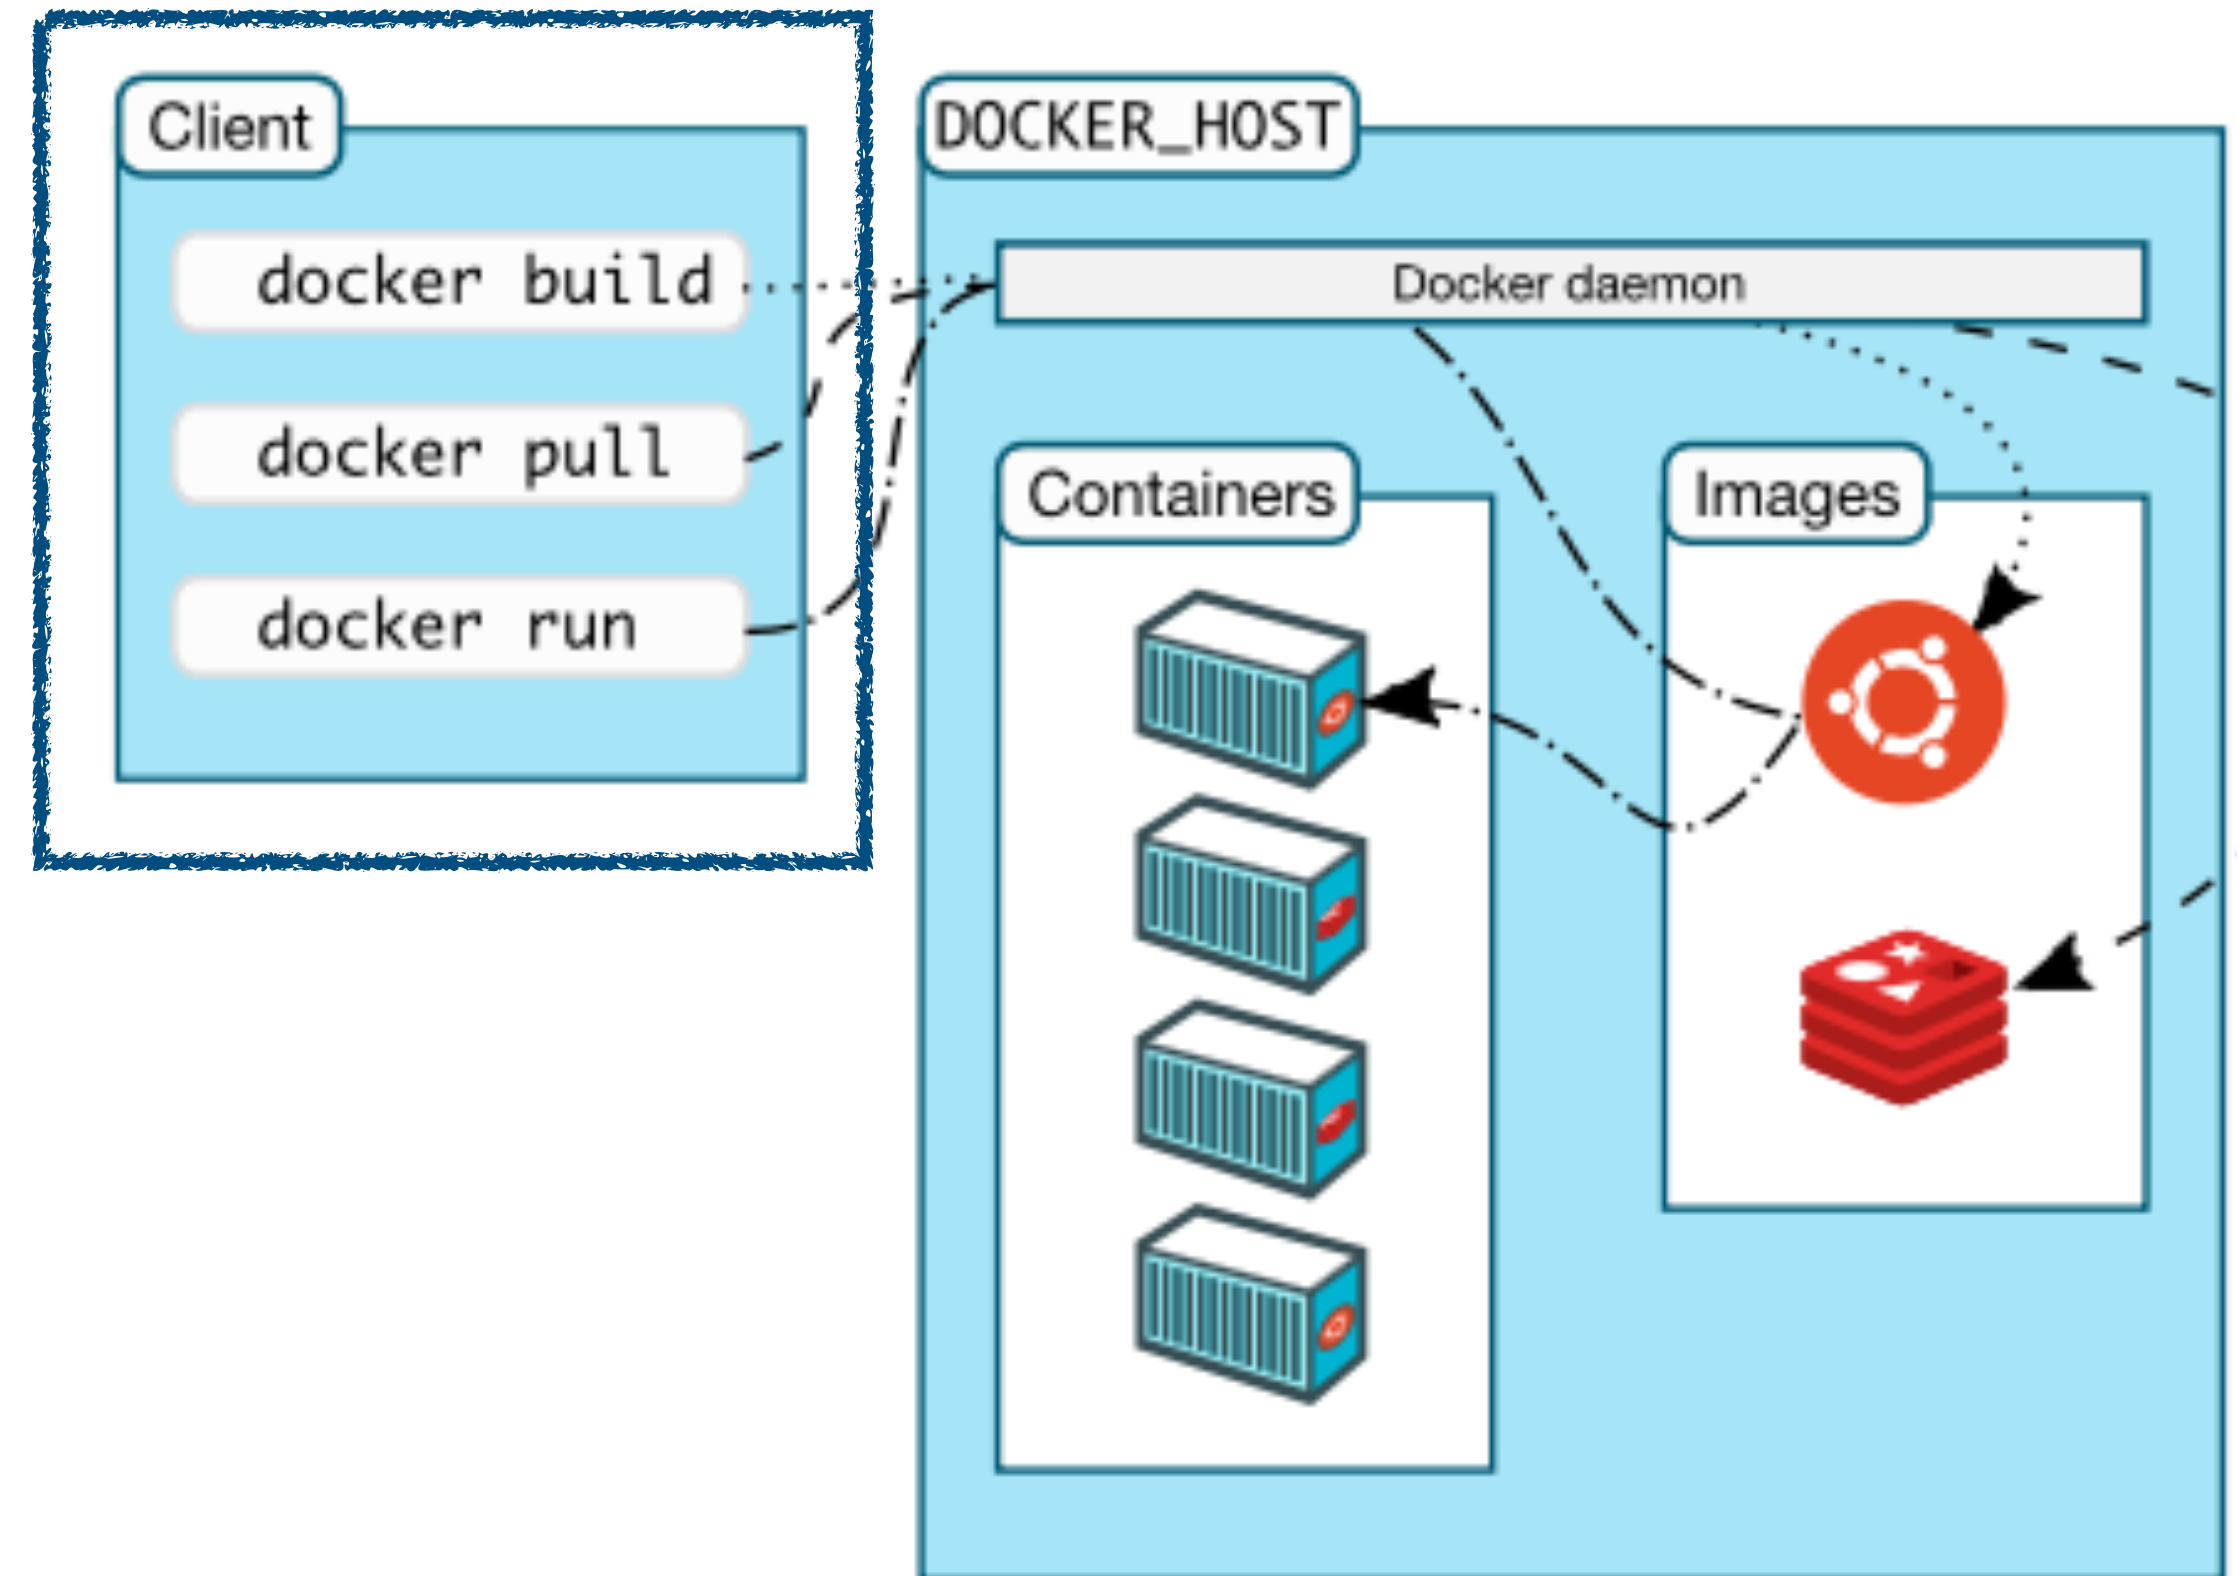
- Most widely-known container platform
- Supports Ubuntu, Fedora, RHEL, CentOS, Windows, etc
- <https://www.docker.com>



Docker

Docker Client

- Component used by users to interact with the Docker platform (daemon)
- Exposes the **Docker API** for:
 - ▶ Running and managing containers
 - ▶ Managing networks and volumes
 - ▶ Reading logs and metrics
 - ▶ Pulling and managing images
 - ▶ ...



Docker

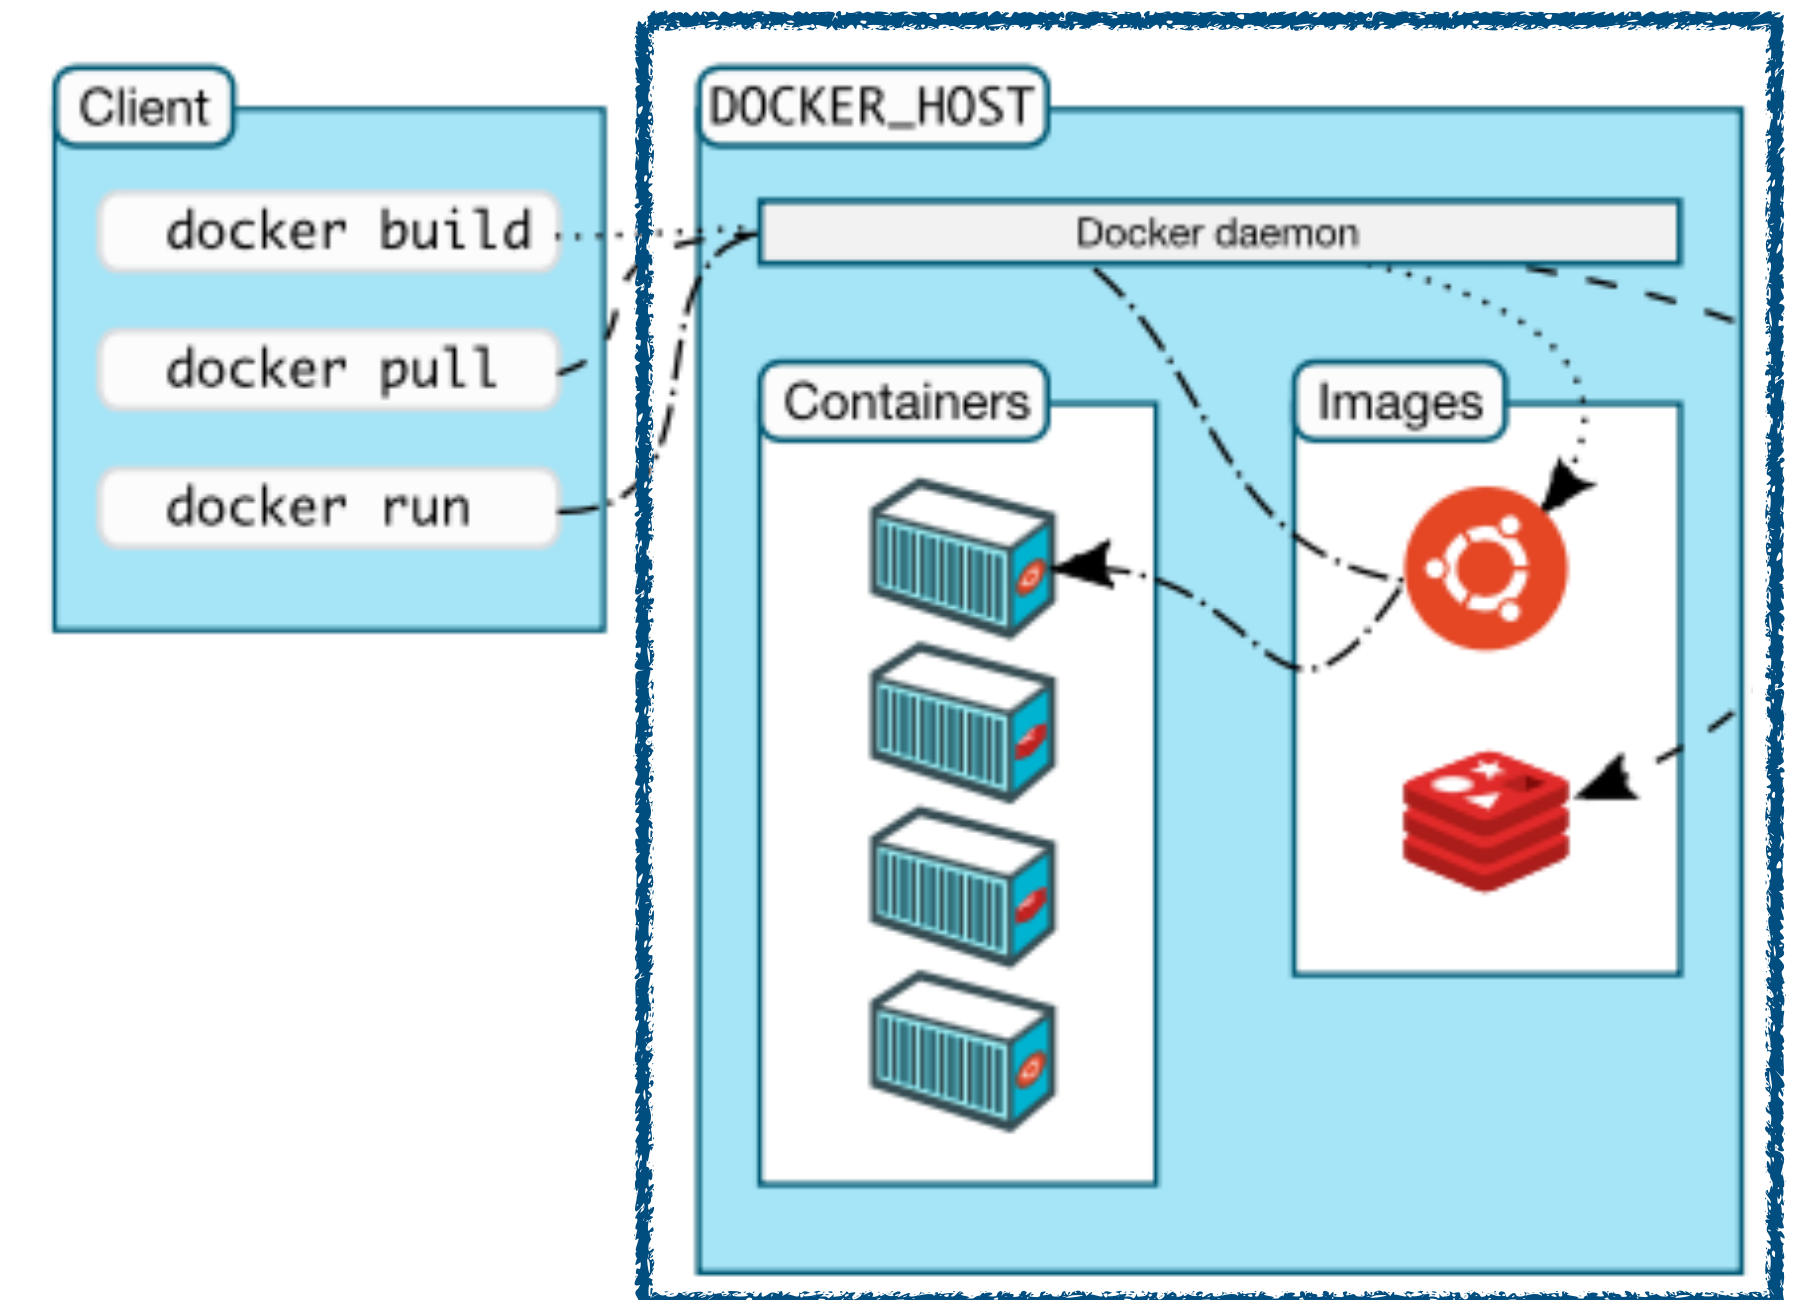
Docker Daemon and Objects

● Docker Daemon

- Uses the Docker API for receiving requests from the Docker Client
- **Manages Docker** Images, Containers, Volumes, Networks

● Docker Objects

- **Image:** immutable (unchangeable) file that contains the source code, libraries, and other files needed for an application to run
- **Container:** runnable instance of an image



Docker

Volumes

- ◎ Containers have an **internal file system** that is **ephemeral** i.e., the container's data is deleted once the container is removed
- ◎ Containers can **mount a file or directory from the host machine**. Stored data is independent from the container's internal file system and **persisted even if the container is removed**
 - **Bind mount:** Generic directory from the host. Any container or host process can access it
 - **Volume:** Special host directory managed by Docker and accessible only by containers

Find more about Docker Volumes at: <https://docs.docker.com/storage/volumes/>

Docker

Network

● **Host:**

- Shares the host networking namespace
- Container services are presented in the network as if run by the host
- Ports are shared (e.g., port 80)

● **Bridge:**

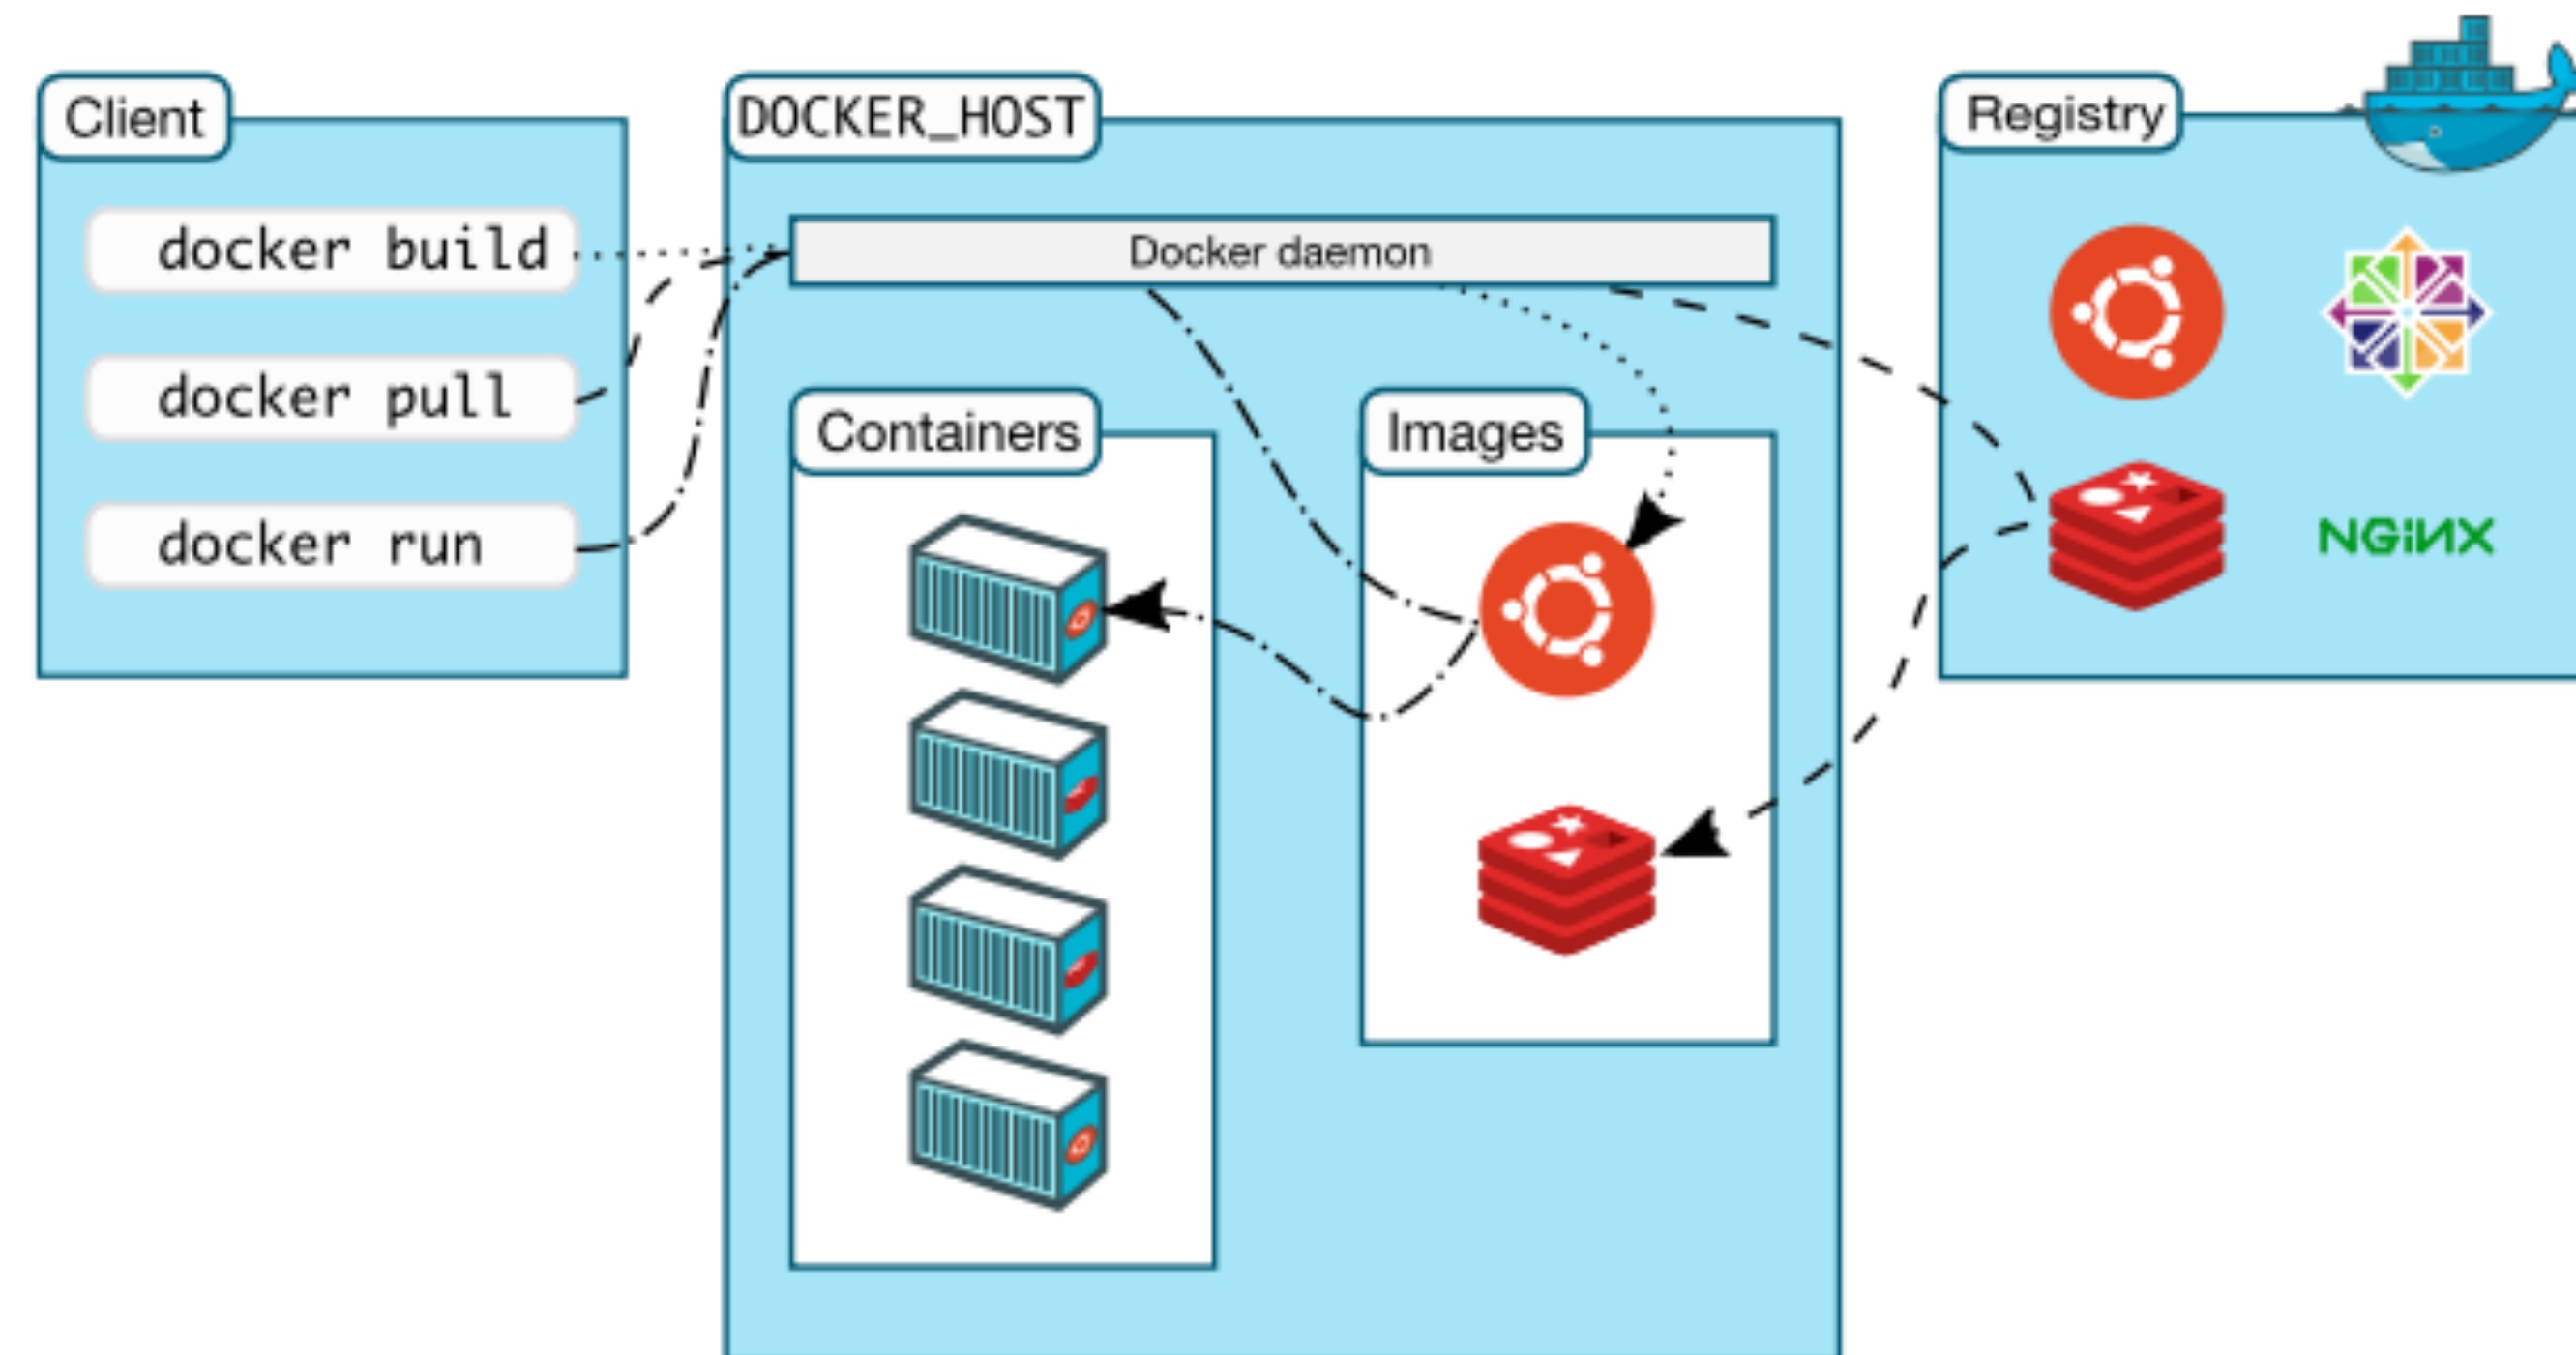
- The container is seen as another node in the physical network

Find more about Docker Networks at: <https://docs.docker.com/network/>

Docker

Docker Registry

- Repository of Docker Images
- Example: <https://store.docker.com>

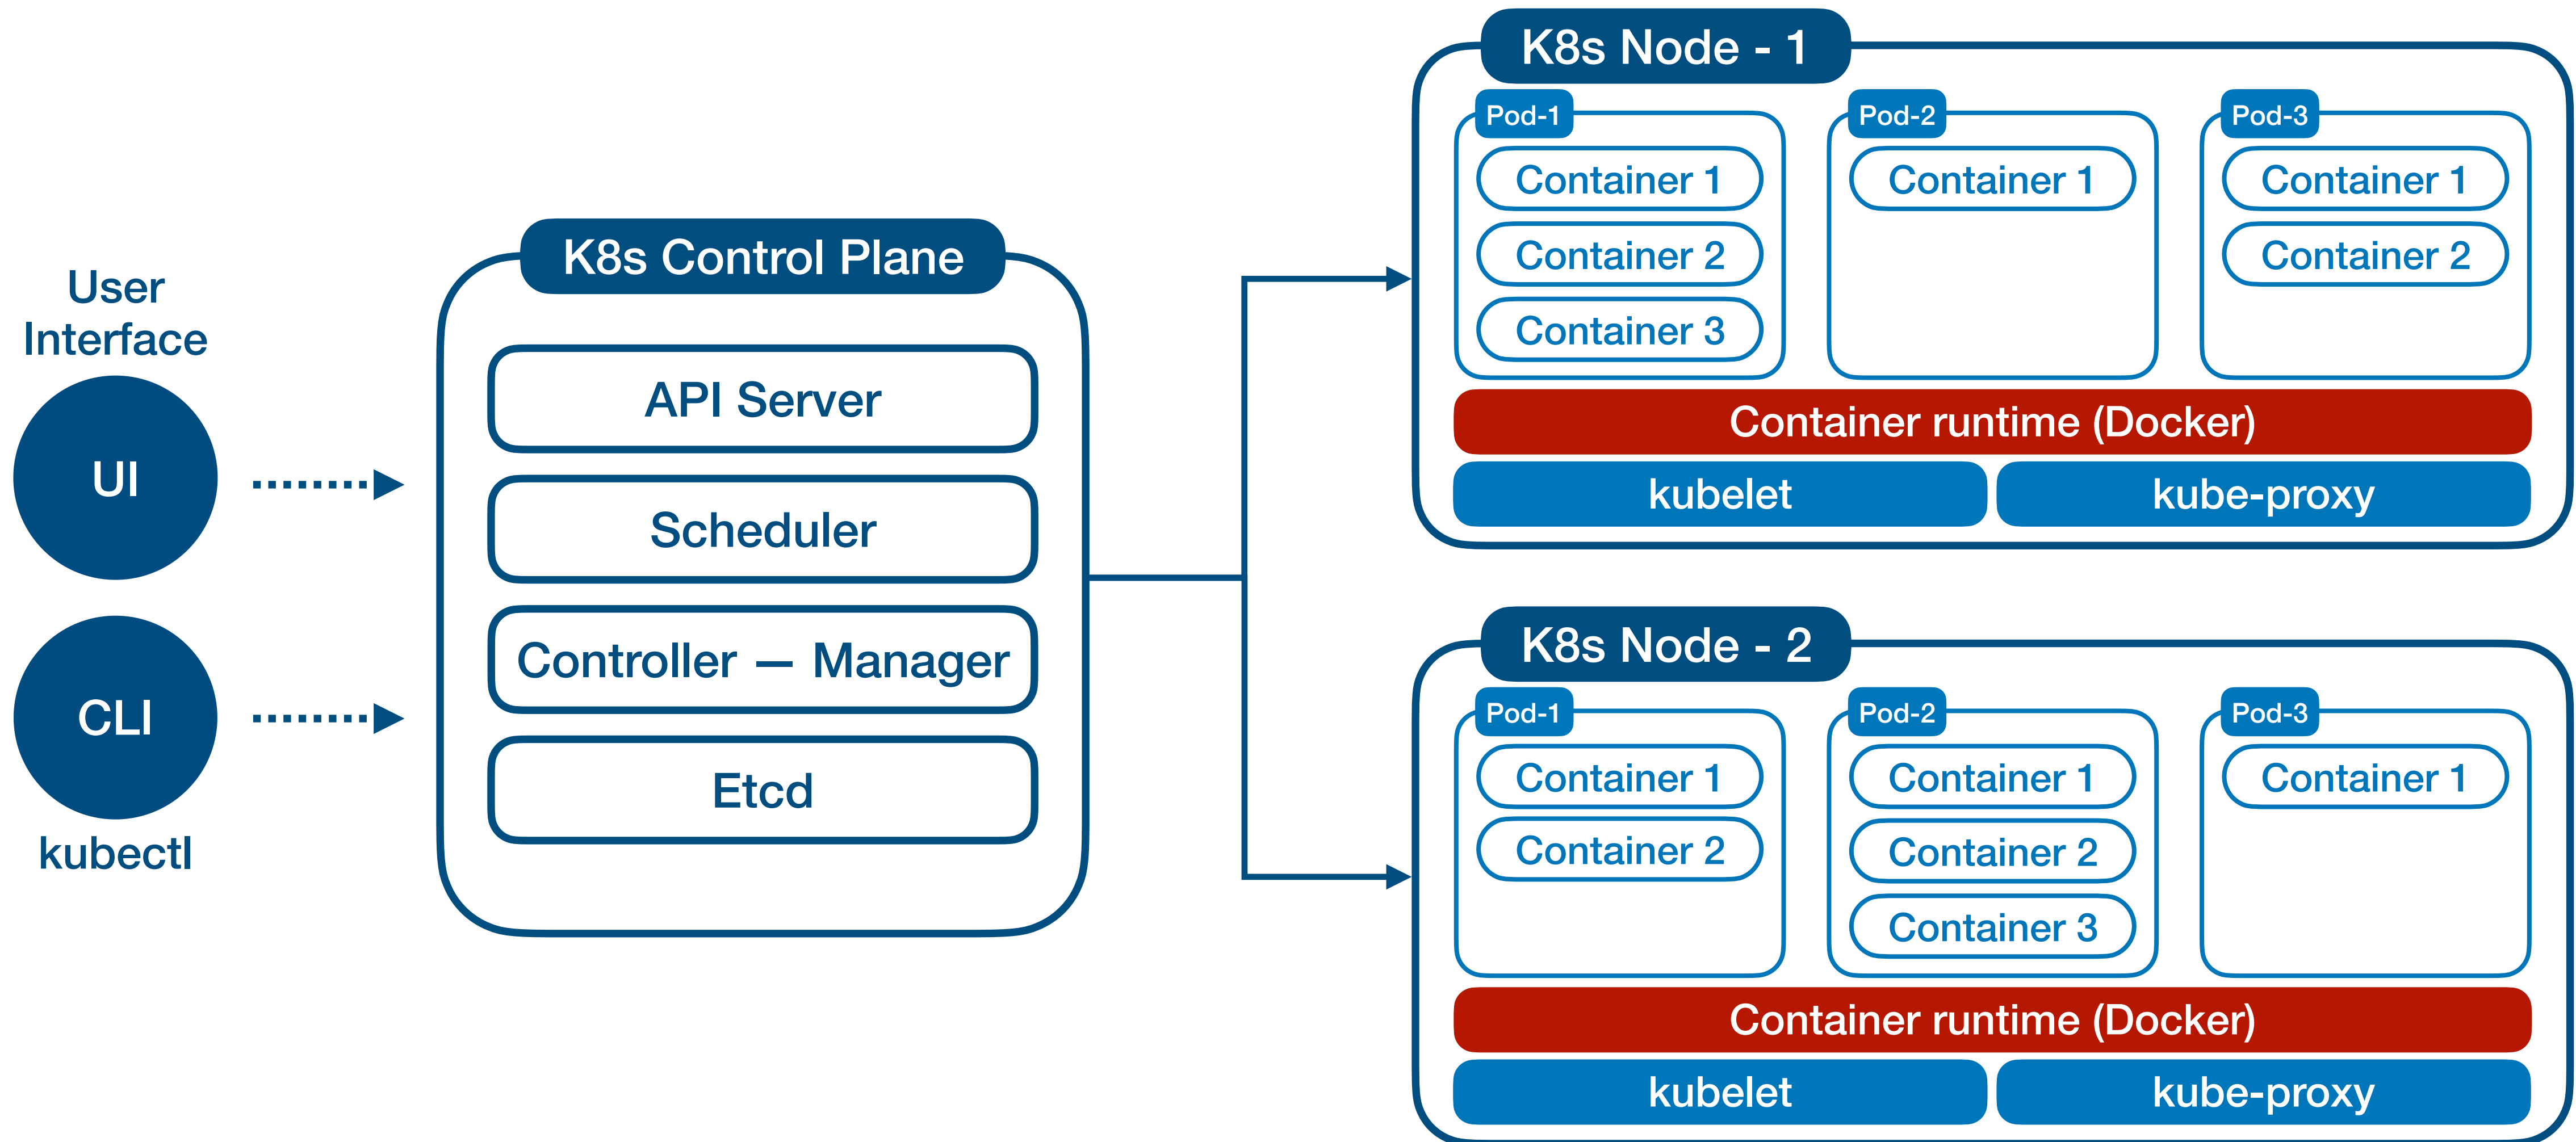


Kubernetes (aka K8s)



- Automates the deployment, scaling and management of containers
- Interesting Features:
 - Network management (e.g., service discovery, load balancing)
 - Modular storage orchestration (e.g., iSCSI, NFS, Ceph, AWS, GCP)
 - Simplified **scheduling, self-healing** and **scale out** for containers
- <https://kubernetes.io>

Kubernetes Components



Kubernetes Components

In brief...

● Control Plane Components

- **API server:** The core component server that exposes the K8s HTTP API
- **etcd:** Key-value store for all cluster configuration data
- **Scheduler:** Distributes unscheduled pods across the available nodes
- **Controller—Manager:** Runs controllers (e.g., for replication, namespaces, ...) to implement K8s API behavior


● Node Components

- **Kubelet:** Manages containers based on incoming Pod specification
- **Kube-proxy:** Accepts and controls network connections to node's Pods
- **Container runtime:** Software responsible for running containers (e.g., Docker)

Kubernetes Cluster

Components

- A K8s cluster is a group of a **Control Plane** and a **set of Nodes**
 - The **Nodes** host Pods that run containerized applications
 - The **Control Plane** manages the Nodes, Network, Storage and Pod resources in the cluster



Kubernetes Cluster

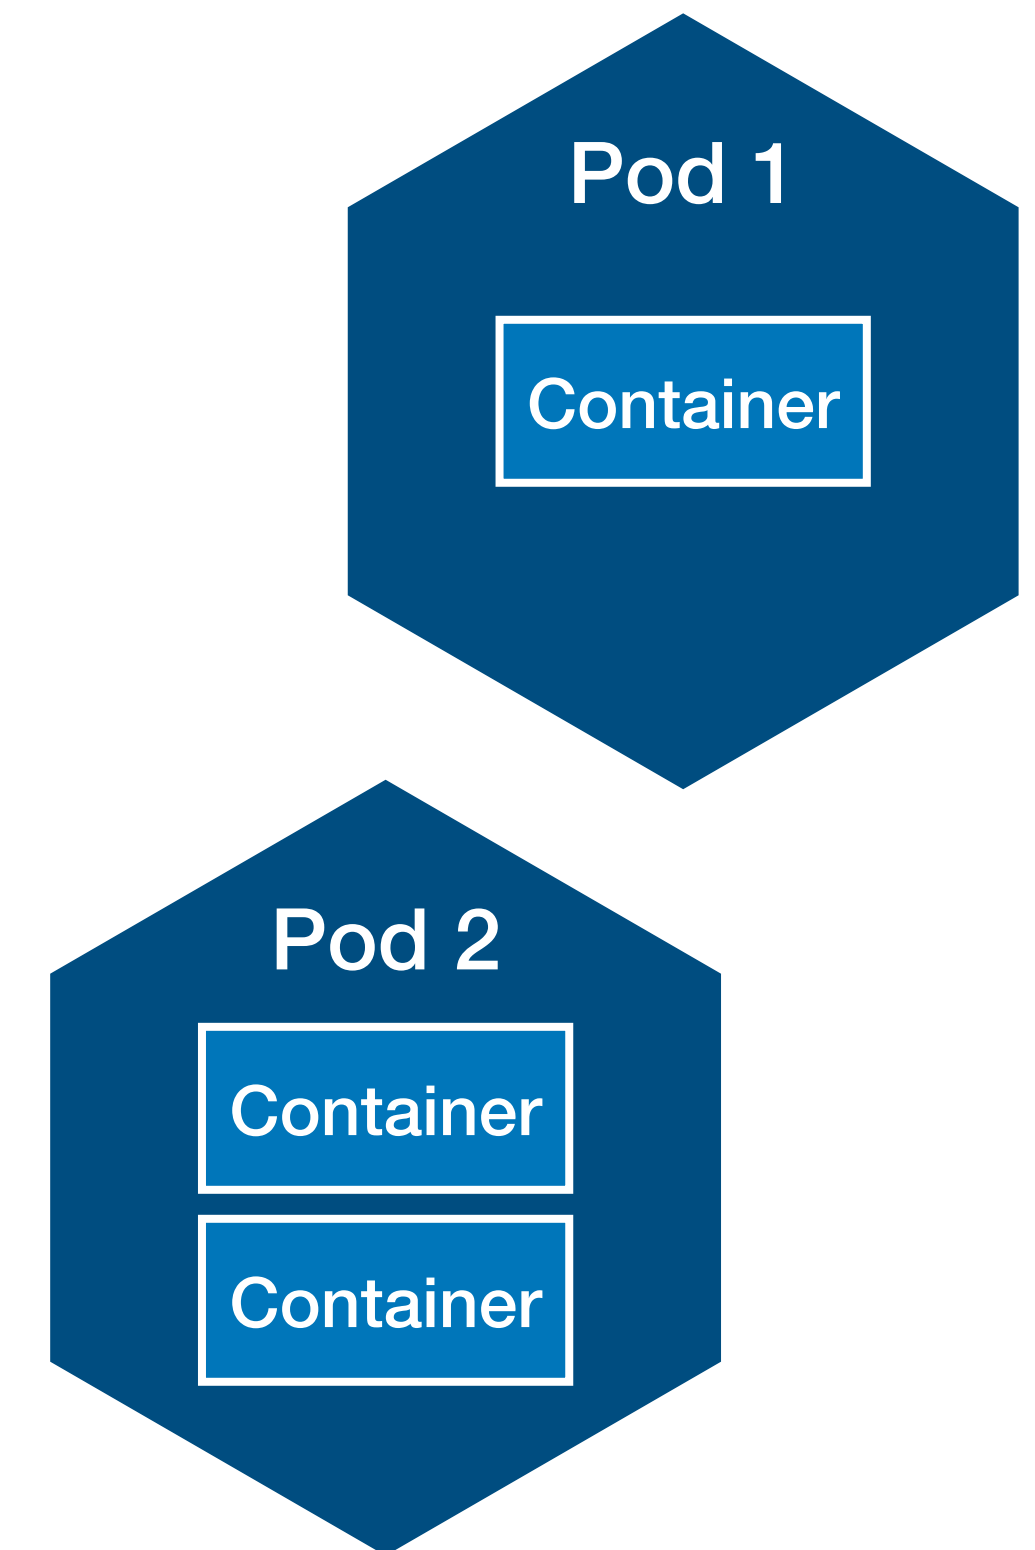
Usage

- An **administrator** uses the **kubeadm** command-line tool to set up a cluster:
 - To initialize a cluster: `kubeadm init`
 - To destroy a cluster: `kubeadm reset`
- **Clients** interact with the cluster through the Control Plane using the **kubectl** command-line tool

Kubernetes Workloads

The Pod computing unit

- A **workload** is an application running on Kubernetes
 - Whether it is a single component or several that work together, on Kubernetes you run these inside a set of **Pods**
- A **Pod** is a group of **one or more containers** with shared storage and network resources and a specification for how to run the containers

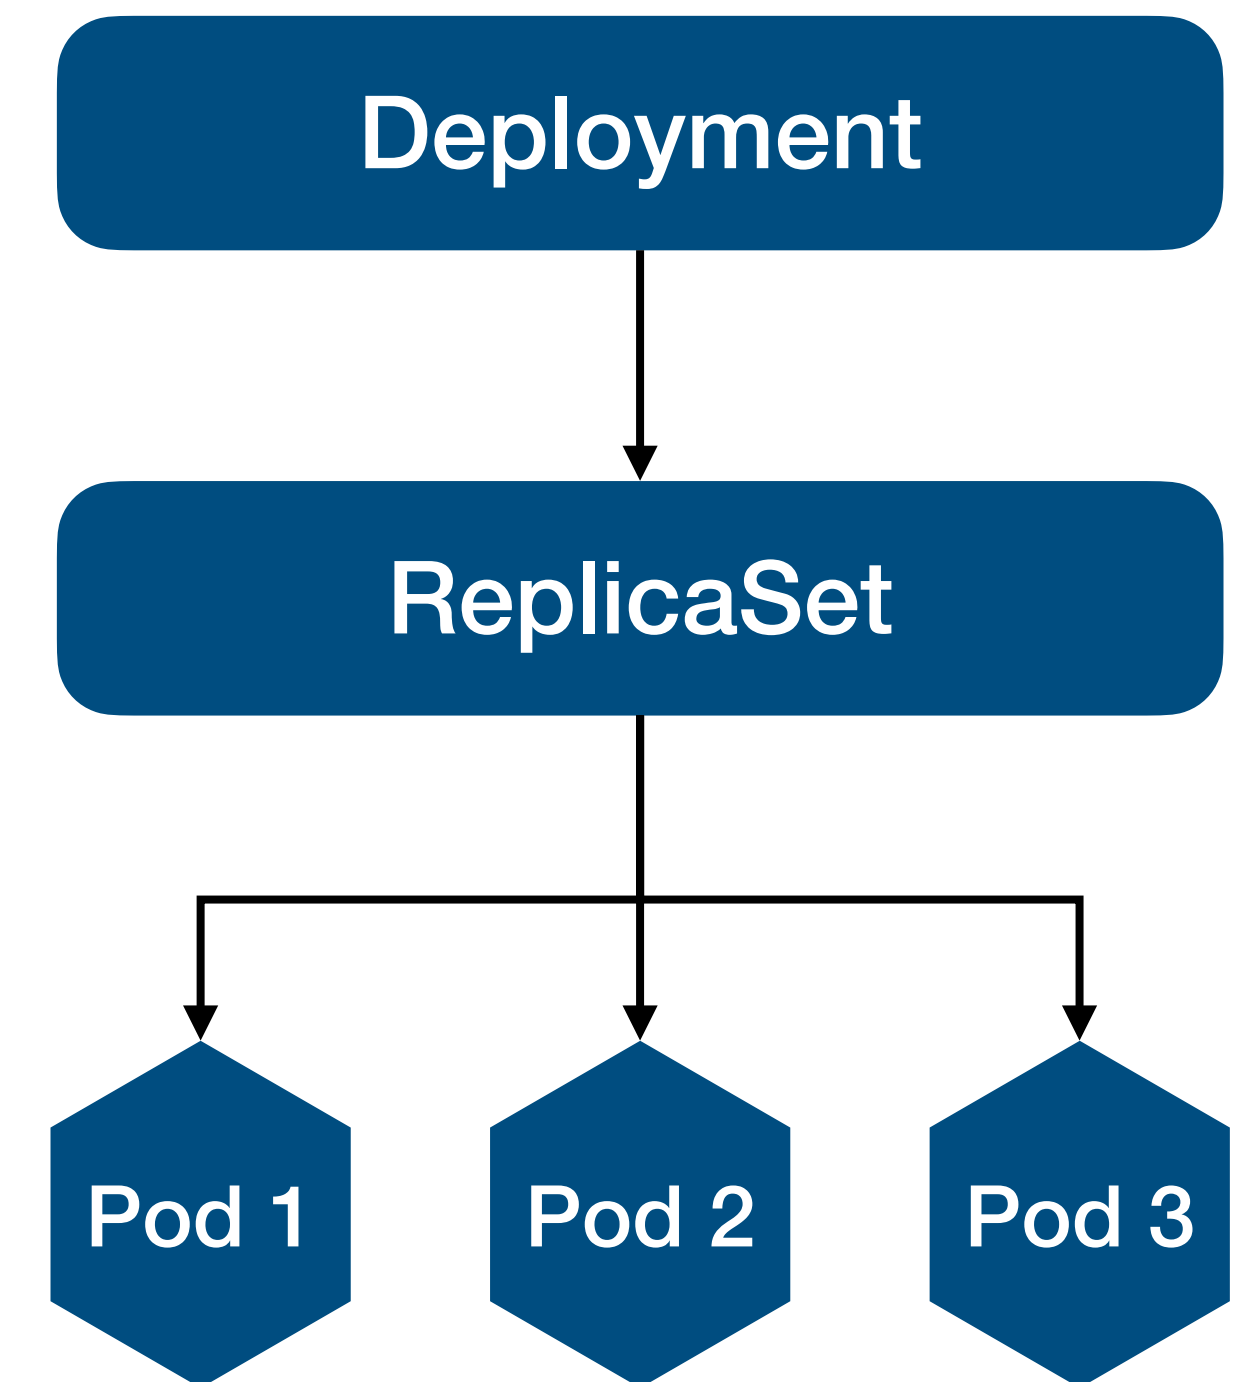


Find more about Pods at: <https://kubernetes.io/docs/concepts/workloads/pods/>

Kubernetes Workloads Resources

Runtime environment for Pods

- **Deployments** provide a declarative way to manage and scale a set of identical Pods (replicas)
 - These define a desired state for the application (e.g., the desired number of replicas) and manage the lifecycle of the corresponding Pods
- **ReplicaSets** ensure that a specified number of Pod replicas are running at any given time
 - The Deployment is a higher-level concept that automatically manages ReplicaSets.




Find more about Workload Resources at: <https://kubernetes.io/docs/concepts/workloads/controllers/>

Kubernetes Network

Network Overlay

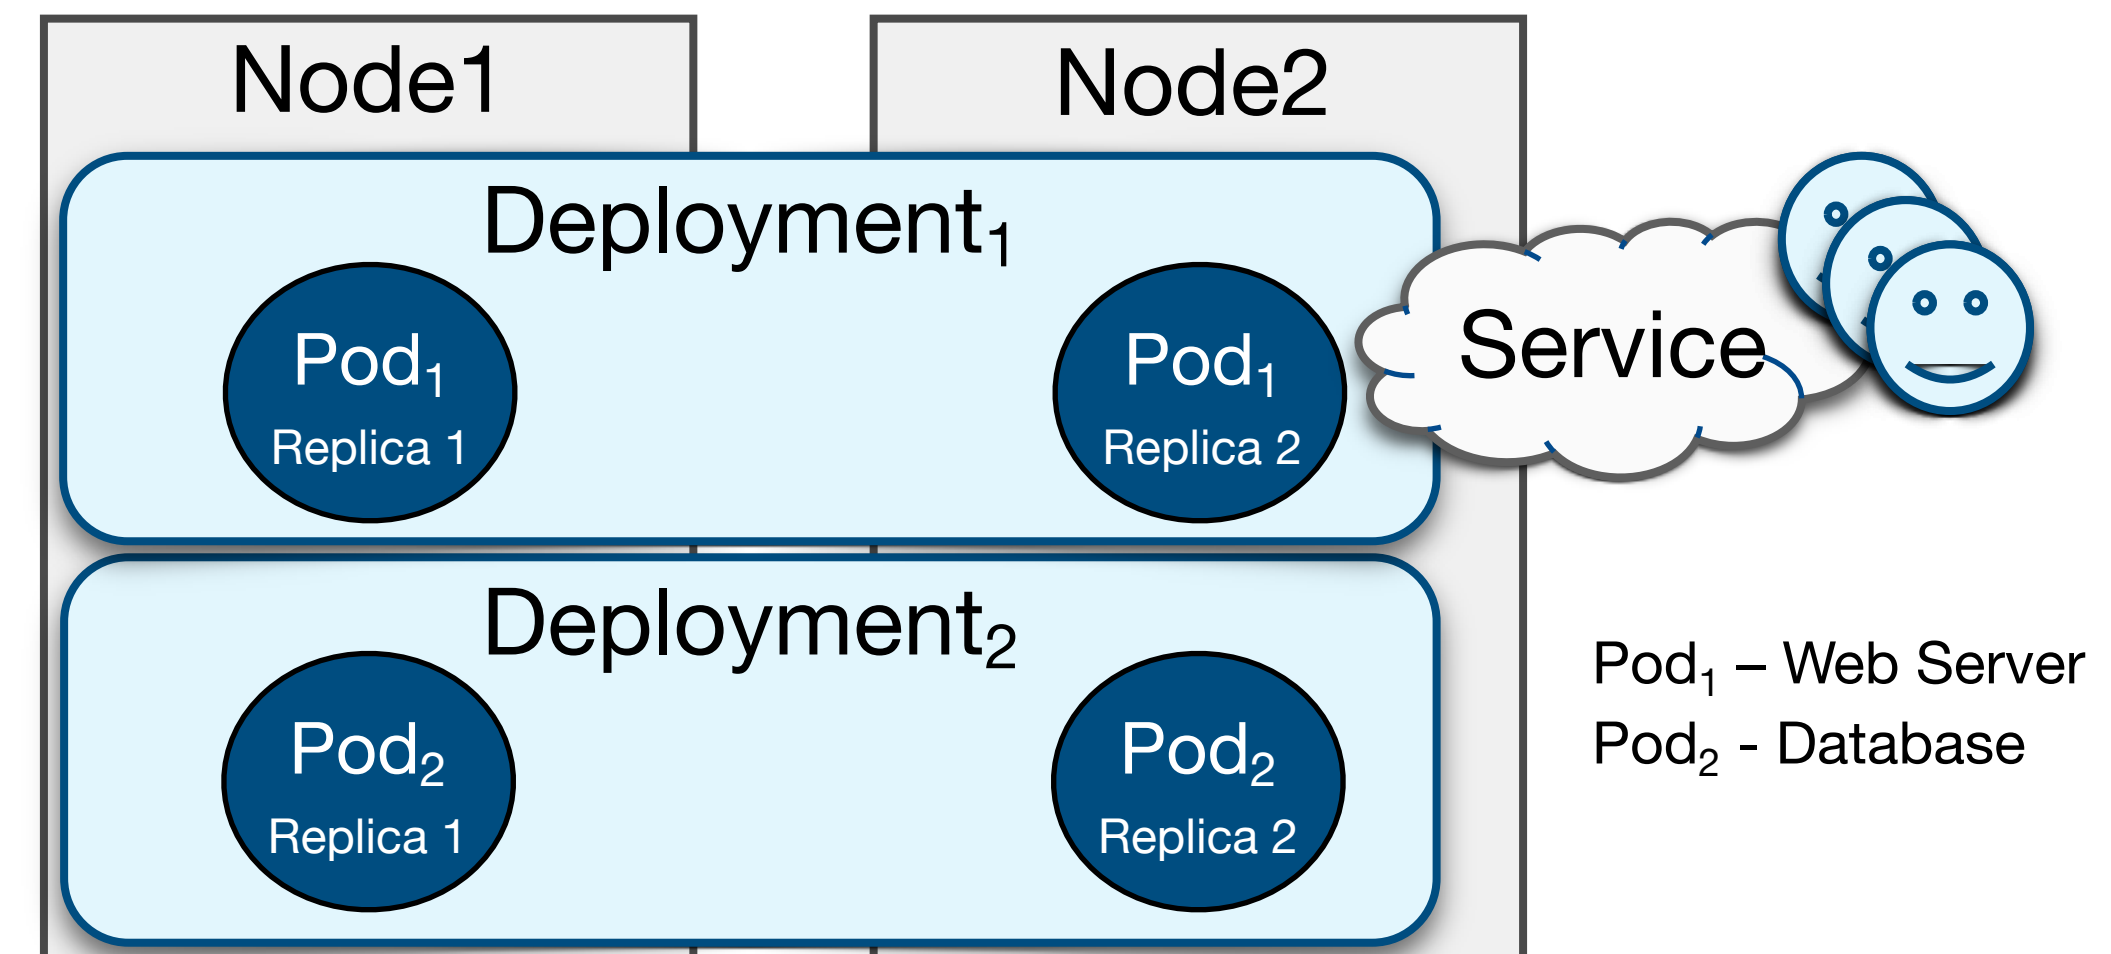
- Each pod has a unique cluster-wide IP address
 - **Networking across pods** (even in distinct worker nodes) is managed through **network overlays** (e.g., Flannel, Calico)



Kubernetes Services

Exposing Pods as a network service

- A **Service** is an abstraction, built on top of the network overlay, for exposing groups of Pods over the network
 - ▶ Each Service object defines a logical set of endpoints (e.g., Pods) and a policy about how to make those pods accessible
 - ▶ Different types of Service policies:
 - **ClusterIP:** Exposes the Service on a cluster-internal IP
 - **NodePort:** Exposes the Service on each Node's IP at a static port
 - **LoadBalancer:** Exposes the Service externally using a external load balancer




Find more about Services at: <https://kubernetes.io/docs/concepts/services-networking/service/>

Kubernetes Storage

Storage Types

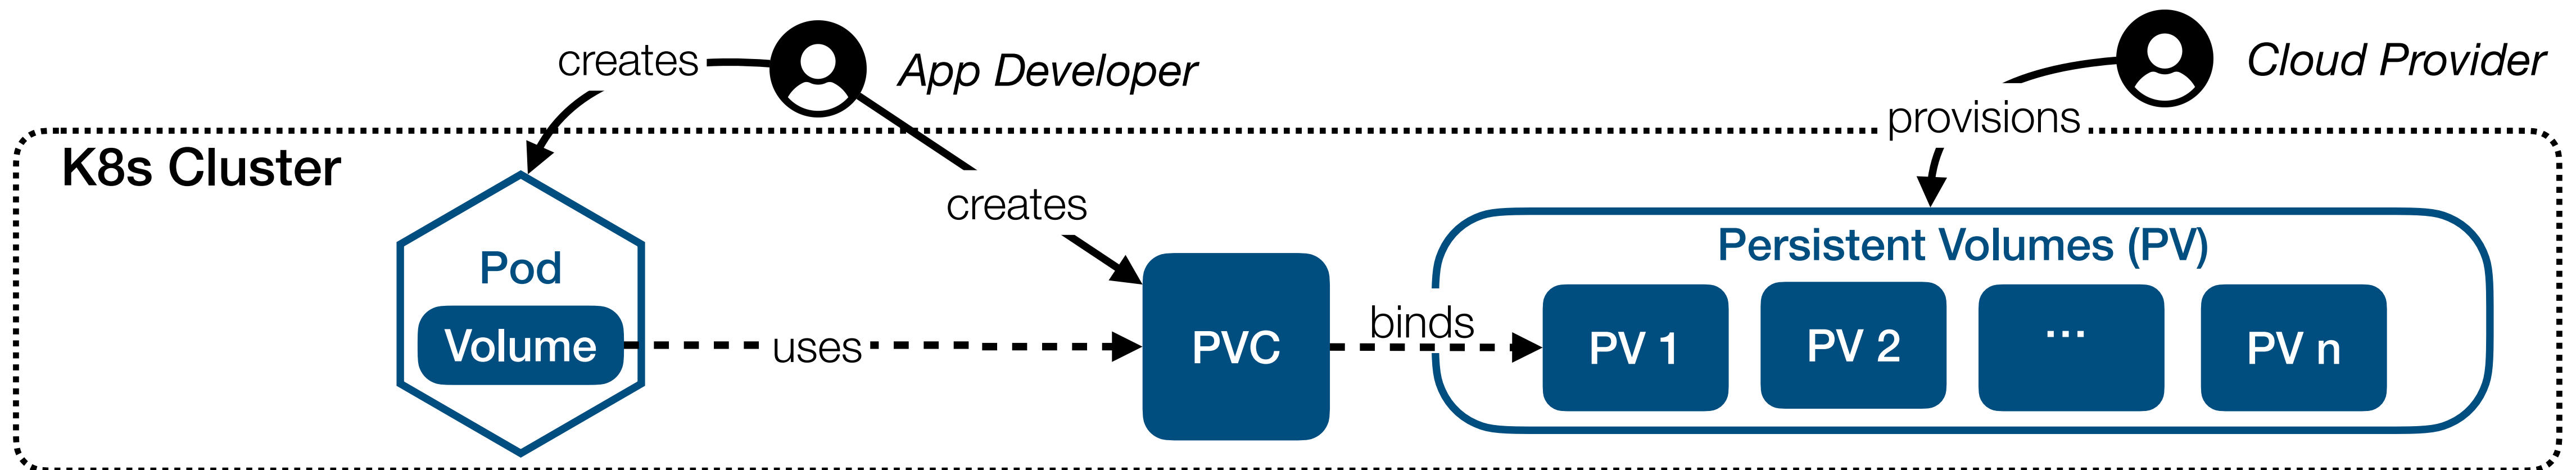
- Pods can access storage volumes provided by different storage backends (*e.g.*, iSCSI, NFS, AWS, GCP)
 - **Ephemeral** storage – available during the pod's lifetime
 - **Persistent** storage – available even if the pod is terminated



Kubernetes Volumes

for Pods

- ◎ **PersistentVolume (PV):** A piece of storage in the cluster provisioned by an administrator or dynamically provisioned using **Storage Classes**
- ◎ **PersistentVolumeClaim (PVC):** A user request of storage for a Pod
- ◎ **Storage Class (SC):** Provides a way for administrators to describe the types of storage they offer (e.g., faster/slower, local/remote storage)



Find more about *Persistent Volumes* at: <https://kubernetes.io/docs/concepts/storage/persistent-volumes/>

Kubernetes (level up!)

Enhancing your Pods configurations and security

● ConfigMaps

- A k8s object for storing and updating **non-confidential** Pod configurations in a key-value pair format (e.g., environmental variables, command-line arguments)

● Secrets

- A K8s object to safely store **confidential** data from Pods (e.g., passwords, tokens, keys)

● Useful for the practical assignment!

Find more at: <https://kubernetes.io/docs/concepts/configuration/configmap/>
and <https://kubernetes.io/docs/concepts/configuration/secret/>

Summary

Containers vs VMs – Disclaimer!

Each technology is built with different goals and the best one depends on the targeted use case!

- VMs are useful when **full server (OS) virtualization** is needed
- Containers are useful for managing virtual environments with heterogeneous libraries and/or applications

Advantages and Disadvantages

Containers vs VMs

● Advantages of Containers over VMs




- **Faster testing/provisioning/migration** (containers are more lightweight!)
- **Better resource utilization and performance** (again, more lightweight...)
- Can easily be **deployed** on both **physical and virtualized servers**
 - Some VM hypervisors provide nested virtualization (i.e., running a VM inside a VM)

● Disadvantages of Containers over VMs

- **Weaker isolation/security** (remember that OS/Kernel are shared)
- **Less flexibility in running different OSs** (e.g., Linux, Windows, BSD, ...)

Further Reading

- S. Alapati. *Modern Linux Administration: How to Become a Cutting-edge Linux Administrator*. O'Reilly, 2016
- K. Morris. *Infrastructure as Code: Managing Servers in the Cloud*. O'Reilly, 2016
- B. Burns. *Kubernetes Up & Running (Second Edition)*. O'Reilly, 2019.



Questions?